

# Lecture 2: Vectors

STAT 450, Eric Fox

## Numeric Vectors

A numeric vector in R is just a collection of numbers. To create a vector use `c()`, which is short for combine.

```
x <- c(1, 2, 3, 47)
y <- c(1, -1, 2, -2)
```

We can add and multiply vectors. We can also multiply or divide a vector by a scalar. Operations are performed element-wise.

```
x + y
```

```
## [1]  2  1  5 45
```

```
x * y
```

```
## [1]  1 -2  6 -94
```

```
2*x # multiplies each element by 2
```

```
## [1]  2  4  6 94
```

```
x/2 # divides each element by 2
```

```
## [1] 0.5 1.0 1.5 23.5
```

```
x^2 # squares each element of the vector
```

```
## [1]  1  4  9 2209
```

## Sequences

Some ways to create vectors containing sequences of numbers:

```
x <- 1:40
x
```

```
## [1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25
## [26] 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
```

```
y <- seq(2, 20, by=2)
y
```

```
## [1] 2 4 6 8 10 12 14 16 18 20
```

```
z <- seq(0, 1, by=0.1)
z
```

```
## [1] 0.0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1.0
```

Use `rep()` to repeat elements in a vector:

```
rep(1, 10) # repeats 1 ten times
```

```
## [1] 1 1 1 1 1 1 1 1 1 1
```

```
x <- c(rep(2,3), rep(3,4))
x
```

```
## [1] 2 2 2 3 3 3 3
```

**Exercise:** Use `seq()` to create a numeric vector containing the sequence 0, 10, 20, ..., 100

```
# solution
x <- seq(0, 100, by=10)
x
```

```
## [1] 0 10 20 30 40 50 60 70 80 90 100
```

## Vector Functions

R provides some easy-to-use functions for computing numerical summaries of vectors. Below are some examples using a vector containing the ages of 10 people.

```
age <- c(54, 40, 53, 26, 56, 52, 28, 39, 15, 17)
length(age) # length of vector
```

```
## [1] 10
```

```
sort(age)
```

```
## [1] 15 17 26 28 39 40 52 53 54 56
```

```
sort(age, decreasing = TRUE)
```

```
## [1] 56 54 53 52 40 39 28 26 17 15
```

```
min(age)
```

```
## [1] 15
```

```
max(age)
```

```
## [1] 56
```

```
mean(age)
```

```
## [1] 38
```

```
median(age)
```

```
## [1] 39.5
```

```
sd(age) # standard deviation
```

```
## [1] 15.70563
```

Or more conveniently, use `summary()` to compute several summary statistics at once.

```
summary(age)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      15.00   26.50   39.50   38.00   52.75   56.00
```

Note that 1st Qu. is the first quartile (25th percentile) and 3rd Qu. is the third quartile (75th percentile).

Documentation on these functions is provided in the help menu. To access the help menu from the console use the `help()` function. For example, enter the following command to read about the `sort` function in the help menu:

```
help(sort)
```

## Data Types

There are four common data types for vectors: numeric (also called double), integer, character, and logical. Some examples:

```
x <- c(1, 0.5) # numeric
x <- 1:10 # integer
x <- c("a", "b", "c") # character
x <- c(TRUE, FALSE) # logical
```

You can use `class()` to check the data type.

```
y <- c(T, T, F, F)
class(y)
```

```
## [1] "logical"
```

```
letters
```

```
## [1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q" "r" "s"
## [20] "t" "u" "v" "w" "x" "y" "z"
```

```
class(letters)
```

```
## [1] "character"
```

Note that for logical vectors T is short for TRUE and F is short for FALSE.

## Coercion

Vectors in R can only contain data of the same type. When different data types are mixed in a vector, R will implicitly coerce all elements to the same type.

```
x <- c("a", 2)
x
```

```
## [1] "a" "2"
```

```
class(x)
```

```
## [1] "character"
```

**Exercise:** Consider the R code below with vectors that contain different data types. Since this is not allowed, what data type does R coerce each vector to?

```
x <- c(1, 0, FALSE, TRUE)
y <- c(1, 2, "three")
z <- c("TRUE", FALSE)
```

```
# solution
```

```
x
```

```
## [1] 1 0 0 1
```

```
class(x)
```

```
## [1] "numeric"
```

```
y
```

```
## [1] "1"      "2"      "three"
```

```
class(y)
```

```
## [1] "character"
```

```
z
```

```
## [1] "TRUE"  "FALSE"
```

```
class(z)
```

```
## [1] "character"
```

## Subsetting a vector

To subset specific elements of a vector use brackets []. The number in the bracket is the index, that is, the position of the element within the vector.

```
age <- c(54, 40, 53, 26, 56, 52, 28, 39, 15, 17)
```

```
age[1]
```

```
## [1] 54
```

```
age[5]
```

```
## [1] 56
```

```
age[1:5]
```

```
## [1] 54 40 53 26 56
```

```
age[c(1, 3)]
```

```
## [1] 54 53
```

```
age[-1] # drop first element

## [1] 40 53 26 56 52 28 39 15 17

age[-c(1, 3)] # drop first and third element

## [1] 40 26 56 52 28 39 15 17
```

Vectors can also be subsetted using logical values.

```
age > 50

## [1] TRUE FALSE TRUE FALSE TRUE TRUE FALSE FALSE FALSE

age[age > 50] # subset ages greater than 50

## [1] 54 53 56 52

sum(age > 50) # counts the number of ages greater than 50

## [1] 4
```

**Exercise:** Test your knowledge of vector subsetting by predicting the output of the following code:

```
weight <- c(140, 139, 187, 181, 131)
weight[2:3]
weight[weight > 150]
sum(weight > 150)
```

## Missing Data

Missing values are denoted as `NA` (not available). The function `is.na()` can be used to check for missing data; it returns a logical vector that is `TRUE` if an element is `NA`, and `FALSE` otherwise.

```
x <- c(7, 11, NA, NA, 12)
is.na(x)
```

```
## [1] FALSE FALSE  TRUE  TRUE FALSE
```

A common task in data analysis is to remove missing values:

```
!is.na(x)
```

```
## [1]  TRUE  TRUE FALSE FALSE  TRUE
```

```
x[!is.na(x)]
```

```
## [1]  7 11 12
```

The `!` symbol is the `not` operator. It negates the elements of a logical vector.